A COMPLEMENT TO COMMITTED ARR FOR THE CONSUMPTION MOTION

Rajesh Kadam · April 2026

Committed ARR captures the signature. We now also need a measure of what happens after it.

Subscription economics rewarded the close. Consumption economics reward the close *and* sustained use. The proposal in this deck is to add a bounded, rules-based second read of the book — not to replace the number the board sees today.

WHERE COMMITTED ARR SERVES US WELL

- Clean signal at the moment of sale: additive, board-recognizable, finance-auditable.
- Integrates with existing comp plans, forecasting, and revenue recognition.
- A stable unit of measurement the board and auditors understand without translation.

WHERE IT UNDER-SERVES THE CONSUMPTION MOTION

- Post-sale adoption is not present in the weekly metric; shelfware, spike-and-drop, and expansion all read identically at signature.
- Field leaders and RevOps lack a triage signal between review cycles.
- Rep coaching is limited to pipeline mechanics, not post-signature outcomes.

This proposal layers a consumption-adjusted read on top of Committed ARR — Committed ARR continues to roll up unchanged

PROPOSAL

cARR — a bounded, rules-based complement to Committed ARR.

A single dollar-denominated figure that preserves additivity and board recognizability, while incorporating trailing-90-day utilization, expansion, shelfware, spike-drop, and new-logo ramp treatment.

$$\text{cARR} = \text{Committed_ARR} \times \text{HealthScore}$$

HealthScore ∈ [0.40, 1.30] — bounded, interpretable, rules-based

BOUNDED

Sensitivity is capped by design.

HealthScore floor 0.40, ceiling 1.30. The impact of any single account on a rep's compensation or forecast is known in advance.

RULES-BASED

Transparent and defensible.

SQL-expressed rules, not a learned model. Every input is traceable to the spec and defensible in a comp dispute or audit.

RAMP-AWARE

New logos are handled explicitly.

Segment-aware blending: Mid-Market 15–60d, Enterprise 30–120d. Adoption lag during onboarding does not flow to the field metric.

ADDITIVE

Committed ARR continues unchanged.

cARR sits alongside Committed ARR on the existing rollup. No change to the figure reported to the board, to finance, or to audit.

Auditable — every rule cites a written spec section

Additive — rolls up by rep, region, segment, and time

Bounded — worst-case account-level impact is known pre-deployment

EVIDENCE FROM THE CURRENT BOOK

cARR on our own data, as of April 2026.

Applied to a representative snapshot of the production book. The figures below are produced by the pipeline, not by the slide — each is reproducible from the SQL models under `pipeline_and_tests/`.

COMMITTED ARR (REPORTED)	CARR (PROPOSED)	IMPLIED ADJUSTMENT	PORTFOLIO HEALTHSCORE
$93.9M	$76.7M	−$17.2M	0.816

Distribution across the 684 in-metric accounts

Healthy	505
At-risk / shelfware	114
Ramping (protected)	58
Spike-drop	5
Expansion	2

What these figures indicate

114 accounts are classified at-risk under the HealthScore rule — reported at full Committed ARR today, with low trailing-90-day utilization.

58 accounts are inside their ramp window. Under v0.5 of the formula they would have taken a HealthScore haircut; D12 ramp blending corrects this before any comp impact.

343 usage logs (~$73K of credits) are excluded as orphans under the D05 rule. Including them would distort each rep's HealthScore denominator.

Live dashboard — rep-level, account-level, and the Ask cARR NL query layer **`gtm-analytics.streamlit.app`**

Source: pipeline_and_tests/data/ · snapshot 2026-04-15

CONSIDERATIONS PRESSURE-TESTED IN DESIGN

Four decisions this proposal has to earn.

Each of these was vetted against the spec and the eval framework before review. The mitigations below are in the pipeline today, not a future roadmap.

Field adoption and compensation trust.

A single-cycle **shadow-comp window** precedes any change in paid compensation. Both figures appear on the monthly rep statement; comp remains paid on Committed ARR. Automated rollback to Committed-based comp if drift exceeds the agreed threshold or any T4 eval fails. The rule is publishable in full — each line of the formula is traceable to a spec section.

Formula integrity and accuracy.

Twelve automated evals across four tiers, from **T1 Correctness** to **T4 Comp safety**. T4 is stop-the-line: a single account where `cARR > Committed` halts deployment. The SQL pipeline is the system of record; the AI layer (Ask cARR) is never authoritative and shows its generated SQL alongside every answer.

Rules-based rather than learned.

Labeled training data for a HealthScore model is not yet available, and a learned multiplier cannot be challenged line-by-line in a comp dispute or audit. Rules-based keeps the model defensible. A learned prioritization layer is proposed for v1.5 — *on top of* the rule-based cARR, not inside it.

New-logo treatment during ramp.

D12 — segment-aware ramp blending. A new Enterprise logo carries the healthy baseline (HealthScore 1.00) from day 1, blending linearly to the measured HealthScore over 30–120 days. Mid-Market blends over 15–60 days. 58 accounts are currently inside that window. Adoption lag during onboarding does not flow to the field metric.

HOW THIS WAS BUILT

Spec-first, with an AI coding assistant in the loop.

Every artifact — metric, pipeline, tests, dashboard — was developed against a written specification with explicit decision records. The assistant accelerated drafting; a human signed off on every merge.

01 · Specifications and decision records

The argument before the code.

Twelve markdown specifications, twelve architecture decision records (ADRs) in a shared decision log. Every code file cites the spec section it implements, and parameter changes require a spec amendment in the same pull request.

- Spec 03 v0.6 — the formula and ramp treatment
- Spec 05 — data-quality catalog (16 assertions)
- Spec 06 — eval framework (T1–T4)
- Spec 11 — the AI product surface

02 · AI-assisted drafting, human judgment

Speed from the assistant, safety from review.

Approximately 70% of boilerplate — SQL scaffolds, test assertions, schema documentation — was drafted by the assistant. The material errors were caught by humans running the code against real sandboxes and reading failing evals.

> **D09:** An assistant-generated MONTH-partitioned BQ schema interacted with the sandbox's 60-day expiration policy, silently evicting 77% of rows on load. Caught at integration, not unit test.

> **D12:** Version 0.5 of the formula produced a uniform HealthScore haircut for every new logo. A dedicated new-logo eval failed and forced the ramp-blend redesign.

Commit-level provenance is preserved: every commit carries a `Co-Authored-By` trailer identifying the assistant and the model.

03 · The AI layer as a product surface

Ask cARR — a bounded NL query layer.

The natural-language query layer is not a thin wrapper around a model. It is a five-stage pipeline, each stage independently testable:

- **Retriever** — mart schema hint
- **Generator** — Claude produces a SQL candidate
- **Verifier** — local EXPLAIN dry-run prior to execution
- **Narrator** — second-pass summary of the result set
- **Refusal policy** — PII, forecasting, cross-rep comp refused explicitly

The generated SQL is shown alongside every answer. The pipeline is the system of record; the model is a query interface.

PHASED ROLLOUT — REVERSIBLE AT EACH GATE

A three-phase rollout with explicit gates.

No phase commits additional compensation liability until the previous phase has produced clean data for one full cycle. Rollback at any stage is a single decision, supported by the pipeline.

PHASE 1

THIS QUARTER

Publish alongside.

cARR joins Committed ARR on the executive dashboard as a second read of the book. Read-only. No compensation impact. Field leaders and reps see both figures side-by-side for one full review cycle.

Gate to Phase 2: under 3% unexplained drift against Committed, all T4 evals green for 90 days.

PHASE 2

NEXT QUARTER

Shadow comp.

Compensation is dual-computed monthly. Both figures appear on the rep statement. *Comp is still paid on Committed ARR.* The objective is to surface divergence between the two and understand its drivers.

Gate to Phase 3: dual-comp delta within the pre-agreed band, no T4 failure, VPS and CFO sign-off.

PHASE 3

Q+2

Cutover with automated rollback.

cARR becomes the driver of variable compensation. Committed ARR continues to report to the board unchanged. Automated rollback to Committed-based comp if drift exceeds 5% or any T4 eval fails — executed by the pipeline, without a manual migration.

Safety net: one-command revert; no data migration required.

RECOMMENDATION — THREE DECISIONS FROM THIS REVIEW

1. **Approve Phase 1** — publish cARR alongside Committed ARR on the executive dashboard this quarter.
2. **Name a VPS delegate and a CFO delegate** to the working group — one point of contact each.
3. **In-principle approval for shadow comp** — contingent on Phase 1 gates being met.

Rollout plan — spec 08 · Rollback gates — spec 08 §4