A NEW GTM NORTH STAR FOR THE CONSUMPTION ERA

Rajesh Kadam · April 2026

We are celebrating bookings that aren't being consumed.

Our North Star is **Committed ARR at signature**. It goes blind after the ink dries — and in a consumption business, the interesting behavior happens after the ink dries.

WHAT THE CFO SEES TODAY

"We reported $93.9M Committed ARR to the board. How much of that is actually being used?"

- $17.2M — 18% of Committed — is at-risk, shelfware, or spike-drop on our own data.
- Reported revenue is decoupling from realized revenue. Audit risk grows every quarter.
- Consumption models punish this faster than subscription did.

WHAT THE VP OF SALES SEES TODAY

"My reps hit 120% of quota on accounts that quietly churn nine months later. I can't coach that."

- One binary signal at signature, then 18 months of silence. Post-sale behavior never reaches the rep.
- Shelfware and expansion look identical in the weekly. Both show "100% of commit signed."
- Compensation rewards the signature, not the outcome.

In a consumption business, ARR without a HealthScore is a vanity metric.

PROPOSAL

One number, still in dollars, that sees past the signature.

cARR preserves the familiar **$** anchor and the additivity the CFO relies on, while adjusting for the post-sale reality the VP of Sales can't see today. It is bounded, rule-based, and auditable.

$$\text{cARR} = \text{Committed_ARR} \times \text{HealthScore}$$

HealthScore ∈ [0.40, 1.30] — bounded, interpretable, no ML

BOUNDED

No paycheck whipsaw.

HealthScore floor 0.40, ceiling 1.30. No single account can swing a rep's number more than −60% / +30%.

READABLE

Reps can read the rule.

Rules-based SQL, not a learned model. A rep can trace every dollar of their comp to a specific line of the formula.

RAMP-AWARE

New logos are protected.

Segment-aware blend: MM gets 15–60d, ENT gets 30–120d of ramp protection. Adoption lag doesn't punish the rep.

BOARD-SAFE

Committed still reports.

cARR *sits alongside* Committed ARR. Nothing the board sees today disappears. We layer truth onto the existing number.

Auditable — every line of the formula traces to a written spec

Additive — rolls up by rep, region, segment

Bounded — paycheck risk is mathematically capped

WHAT CARR REVEALS — ON OUR OWN BOOK, AS OF APRIL 2026

The 18% of Committed ARR that isn't actually there.

Running the proposed metric over the current snapshot surfaces the accounts where commit and consumption have silently decoupled. These numbers come from the pipeline, not a slide template.

COMMITTED ARR	CARR	DELTA	WEIGHTED HEALTHSCORE
$93.9M	$76.7M	−$17.2M	0.816

Where the 684 in-metric accounts land

Healthy	505
At-risk / shelfware	114
Ramping (protected)	58
Spike-drop	5
Expansion	2

Why this matters today

114 at-risk accounts are being reported at full Committed ARR on today's board deck. They are the 18% gap.

58 ramping accounts would have been unfairly penalized in v0.5 of the formula. Ramp protection (D12) fixes that before a rep ever sees a comp impact.

343 usage logs ($73K) were excluded as orphans. Without the pipeline's D05 rule, they would inflate the denominator and distort every rep's HealthScore.

Live dashboard — rep-level, account-level, and the "Ask cARR" NL query layer `gtm-analytics.streamlit.app`

Source: pipeline_and_tests/data — snapshot 2026-04-18

ANTICIPATED OBJECTIONS

The four questions you are about to ask.

Overhauling how a sales engine is measured and paid is a reversible-cost decision only if we anticipate the pushback. Here's the plan for each.

Q "Will my reps revolt when their comp changes?"

A They won't, because comp doesn't change blind. We run **shadow comp for one full cycle** — paychecks are computed both ways and reps see both numbers every month. If drift exceeds threshold or any T4 eval fails, we **roll back to Committed-based comp automatically**. The rule is publishable; a rep can trace every dollar of their comp to a line of the formula.

Q "What if the formula is wrong?"

A We have **12 automated evals across four tiers**, from T1 Correctness to **T4 Comp safety**. T4 is stop-the-line: a single account where `cARR > Committed` blocks a deploy. The SQL pipeline is the system of record, never the LLM. Every parameter cites its section in spec 03.

Q "Why not use machine learning here?"

A No labeled training data yet, and reps cannot challenge a black-box model in a comp dispute. Rules-based is the **audit story**: every input is defensible. v1.5 layers a learned prioritization *on top of* the rules — never underneath. The LLM in "Ask cARR" is a query interface, not a source of truth.

Q "What about my new logos in their ramp period?"

A **D12 — segment-aware ramp blend.** A new Enterprise logo inherits a healthy HealthScore (1.00) from day 1, fading linearly to the measured HS over 30–120 days. Mid-Market ramps over 15–60 days. **58 accounts** on today's book are inside that window. A rep is never punished for adoption lag they don't control.

HOW THIS WAS BUILT

AI-first, spec-driven. The methodology compounds.

Every artifact in this repo — specs, SQL, tests, dashboard — was built in a spec-first loop with an AI coding assistant. The discipline is what makes it defensible, not the speed.

`01 · Specs first`

Write the argument before the code.

12 markdown specs, 12 ADRs in a decision log, every code file cites the spec section it implements. The panel can open `specs/` and read the argument *before* reading a line of SQL.

- Spec 03 v0.6 — the formula + ramp protection
- Spec 05 — data-quality catalog (16 assertions)
- Spec 06 — eval framework (T1–T4)
- Spec 11 — the AI product surface

`02 · AI drafts, human judges`

The AI is fast but not safe alone.

~70% of boilerplate came from the assistant. The subtle bugs came from humans running the code against real sandboxes and reading failing evals.

D09 caught by humans: AI generated MONTH-partitioned BQ tables; the sandbox's 60-day expiration silently evicted 77% of rows. Found on an integration run, not a unit test.

D12 caught by humans: v0.5 of the formula penalized every new logo. A new-logo eval failed, which forced the ramp-blend design.

Every commit carries a `Co-Authored-By: Claude` trailer. Full provenance.

`03 · AI as a product surface`

"Ask cARR" is a productized example.

The NL query layer is not a wrapper around an LLM. It is a five-stage pipeline:

- **Retriever** — schema hint from the mart
- **Generator** — Claude writes SQL
- **Verifier** — local EXPLAIN dry-run before execution
- **Narrator** — second Claude pass summarizes rows
- **Refusal policy** — PII, forecasting, cross-rep comp all refused

The SQL is always shown alongside the answer. The LLM is never the system of record.

PHASED ROLLOUT — REVERSIBLE AT EVERY STEP

Approve Phase 1 today. Comp change waits for evidence.

Each phase has a gate. No phase commits comp liability until the previous phase has produced clean data for a full cycle. Rollback is one decision, not a project.

PHASE 1

THIS QUARTER

Publish alongside.

cARR joins Committed ARR on the exec dashboard. Read-only. No comp impact. Reps and leaders see both numbers side-by-side for one full cycle.

Gate to Phase 2: < 3% unexplained drift, all T4 evals green for 90 days.

PHASE 2

NEXT QUARTER

Shadow comp.

Paychecks are computed both ways every month. Reps see both numbers on their comp statement. *Comp is still paid on Committed*. We learn where the two numbers diverge and why.

Gate to Phase 3: < 2% unexplained paycheck delta, no T4 failure, VPS + CFO sign-off.

PHASE 3

Q+2

Cutover with rollback.

cARR becomes the comp driver. Committed stays on the board. Automatic rollback if drift > 5% or T4 fails — the pipeline flips back to Committed-based comp without a PR.

Safety net: one-command revert; no manual data migration.

THE ASK — THREE DECISIONS TODAY

1. **Approve Phase 1** — publish cARR alongside Committed on the exec dashboard this quarter.
2. **Name a VPS-delegate and a CFO-delegate** to the working group — one point of contact each.
3. **Green-light shadow comp for Phase 2** — contingent on Phase 1 gates.